

Presentations Details

Materials Due 11:59p on Monday, December 2, 2024 to the course website.
Presentations will be on Tuesday, December 3 and Thursday, December 5 in class.

Presenting Regression Results

You will have the opportunity to run a regression and present the results using data of your own choosing. Complete the following steps to prepare a 10-minute presentation of your results. To receive credit for this portion of the final, submit the dataset you used, your R code, and the deck of 5 slides to the course website by **11:59pm on Monday, December 2**. *Pro tip: don't underestimate how long data importing and cleaning can take.*

First, find data and clean it:

1. Find a dataset that contains three variables that you would like to work with. You'll use these variables to assess whether some outcome variable Y can be explained in part by independent variables X_1 and X_2 ; specifically, whether X_1 is related to Y controlling for X_2 . The dataset should satisfy the following requirements: (a) The variable that you'll use as the dependent variable is an interval/ ratio variable that can take at least three values (is not binary). (b) The data are cross-sectional (so do not contain the same units measured at different points in time). You are welcome to make use of a subset of the dataset that satisfies these conditions, for instance, if the data are time series cross-sectional, you can take one cross section and use that.
2. Load the data into R and make a dataframe called `mydata` that contains only the three variables you plan to use.
3. Clean `mydata` so that you can run a meaningful regression. For instance, do the data add extra numbers, such as 99s instead of NAs, that will throw things off if you don't remove them? Is R correctly treating the numeric variables as numbers, or do you need to force R to do so? Are any variables stored as strings that you should convert to numerics?
4. Now, your interval/ratio variable is your dependent variable (Y), which leaves two independent variables (X_1, X_2). Pick one of the two independent variables to be a dummy variable. Recode it so that it takes the value of 1 for some value(s) of the variable and 0 otherwise. This variable will be X_2 . For instance, if it measured age in years, convert it to be coded as 1 if age is above some cutoff number of years and 0 otherwise. (If the variable is already a dummy variable, then leave it and skip to the next step).

Now, analyze the data in R:

5. Regress Y on X_1 and X_2 .
6. Make a well-formatted scatterplot that shows the relationship between X_1 and Y , and that uses 2 different symbols and colors to contrast points for which $X_2 = 1$ and $X_2 = 0$. Include a legend.
7. Add the two regression lines associated with each of the two values of X_2 to the X_1, Y scatterplot.
8. Note the coefficients, their levels of statistical significance, and the 95% confidence intervals from the regression— you'll use this information below.

Now, build 5 slides to accompany your presentation of the results:

9. Make a slide that will aid your narration of why we might expect a causal relationship between X_1 and Y .
10. Make a slide that will aid your narration of where the data came from and what the three variables are.
11. Make a slide that contains your scatterplot with regression lines.
12. Make a slide that contains both your regression equation as well as a table with the estimated regression coefficients, standard errors, p-values, and 95% confidence intervals for both X_1 and X_2 .
13. Make a slide that will aid your narration of why, even if the results were perfectly statistically significant, we should be skeptical of a causal interpretation.

Finally, prepare a 10-minute presentation of your slides:

14. Prepare a presentation to give with your slides **that lasts no longer than 10 minutes** and covers the following:
 - (a) Very briefly introduces a hypothesis for why the concept measured by X_1 could be causally related to the concept measured by Y .
 - (b) Introduces the data and three variables you'll use.
 - (c) Briefly introduces the scatterplot and what it suggests about relationships in the data.
 - (d) Explains the regression run and introduces the regression lines in terms of sign and magnitude.
 - (e) Explains what each regression line tells us and what the difference between them tells us.
 - (f) Explains in what sense the regression "controls for" X_2 .
 - (g) Discusses the statistical significance of the results and what we should make of the hypothesis based on all of this.
 - (h) Explains why regression results that perfectly supported the hypothesis would not *confirm* a causal relationship. (Think about confounders, endogeneity, etc.)
 - (i) Offers one concrete way that future research could take a step toward better assessing this hypothesis. (new data? new analyses?)